

Data C100 - Final

Summer 2026

Write your name BIG and clearly: _____

Email: _____@berkeley.edu

Student ID: _____

Name and SID of the person on your left: _____

Name and SID of the person on your right: _____

Exam Room: _____ Seat Number: _____

Instructions:

This exam consists of **71 points** spread out over **9 questions**. The exam must be completed in **minutes** unless you have accommodations supported by a DSP letter.

- Note that some questions have circular bubbles to select a choice. Please shade in the circle fully to mark your answer.
- **Write clearly and legibly.** We reserve the right to withhold points from answers that are very difficult to read.
- Blank answers and incorrect answers are graded identically, so it's in your best interest to answer every question.
- **You MUST write your Student ID number at the top of each page.**
- You should not use a calculator, scratch paper, or notes you own other than the reference sheets distributed at the beginning of the exam.

For all Python questions, you may assume Pandas has been imported as `pd`, NumPy as `np`, the Python RegEx library as `re`, `matplotlib.pyplot` as `plt`, and `seaborn` as `sns`.

Honor Code:

As a member of the UC Berkeley community, I act with honesty, integrity, and respect for others. I am the person whose name is on the exam, and I completed this exam in accordance with the Honor Code.

Signature: _____

This page has been intentionally left blank.

1 Potpourri [8 Pts]

- (a) [2 Pts] Table A has 200 rows. Table B has 300 rows. The ID column of A has 200 unique values. The ID column of table B has 300 unique values. The tables are joined on the ID column.

Complete the table below by writing the minimum and maximum number of rows that could result from each join.

Join Type	Minimum Number of Rows	Maximum Number of Rows
Inner join of A and B		
Left join of A to B		
Left join of B to A		
Full join of A and B		

- (b) [2 Pts] Let `df` be the following DataFrame:

a	b
10	10
40	30
30	20
20	5
5	5

`z` is a boolean Series with the values `[True, False, False, True, False]`

Write the output of the following. If an answer results in a series, write the values of the series in the order they appear.

Code	Output
<code>np.sum(df[z]['a'])</code>	
<code>np.mean(z)</code>	
<code>np.sum(df[2:3]['b'])</code>	
<code>np.sum(df['a'] == df['b'])</code>	

(c) [4 Pts] For each combination of variables below, write three plot types that could be used to visualize the variables in a single plot.

- You must use plot types that we've taught in Data 100.
- If you write more than three plot types for a variable, you will lose credit.
- If there are only two possible answers for a question, write NA as the third plot type.

Variable type(s)	Plot Type 1	Plot Type 2	Plot Type 3
One quantitative			
Two quantitative			
One quantitative and one qualitative nominal			
Two quantitative and one qualitative ordinal			


2 EDA [10 Pts]

A company has three stores: Store A, Store B, and Store C. Every year, the company identifies the 10 customers with the highest total **count** of transactions combined across all its stores in that year.

- For example, in 2024, suppose a customer makes 3 transactions at Store A, 2 transactions at Store B, and 4 transactions at store C. The company would compare this customer's 9 total transactions against the total number of transactions of all other customers in 2024.

Then, the company computes the total **dollar amount** that the top 10 customers spent at each of the company's stores in the year the customers were selected.

Consider the following visualization of this data:



The visualization above was generated with a single call to `sns.scatterplot` and a single DataFrame called `plot_df`.

(a) [1 Pt] How many rows are in `plot_df`?

(b) [1 Pt] What is the minimum number of columns in `plot_df`?

(c) [8 Pts] `plot_df` was generated by transforming a different DataFrame called `raw_df`. Each row of `raw_df` represents one transaction. Here are the columns of `raw_df`:

- `transaction_id`: a unique identifier for the transaction
- `customer_id`: an identifier for the customer associated with the transaction
- `store_id`: an identifier for the store where the transaction occurred
- `year`: the year of the transaction
- `amount`: the dollar amount of the transaction

Write an analysis plan to transform `raw_df` into `plot_df`. Things to keep in mind:

- The same customer might be in the top 10 for multiple years.
- You can assume the company had at least 10 unique customers each year.
- `plot_df` must only include rows that are needed for the plot. No extra rows are allowed.
- `plot_df` may optionally include columns that are not required to make the plot, since the call to `sns.scatterplot` could ignore these columns.
- Make sure to use the appropriate canonical English terms we learned in lecture. See the reference sheet for examples of these terms.

Write your answer on the next page! →

Repeated descriptions: Every year, the company identifies the 10 customers with the highest total **count** of transactions combined across all of its stores in that year. Then, the company computes the total **dollar amount** that the top 10 customers spent at each of the company's stores in the year the customers were selected.

Reminders: `raw_df` contains transaction ID, customer ID, store ID, year, and amount. The plot shows year on the x-axis and total amount spent by the top 10 customers on the y-axis. The point shape corresponds to the store. You are writing steps to produce `plot_df`.

Hint: You may find it helpful to produce multiple *DataFrames* and then join them. Our solution produces two *DataFrames* and joins them.

To be eligible for partial credit, keep your answer organized and neat.

3 I still love all of my regexes [4 Pts]

Hazel obtains text from the "Staff" page of the Data 100 website. She wants to use RegEx to extract certain pieces of information.

Here is an example string. The character "_" represents a single space, and you may use it in your response.

```
"Name:_Oski;_Age:_161;_Courses_taken:_D100,
_DATA143,_CS61B,_Math54,_Data27s;_Phone Number:
_510-642-9999;_Likes:_[Bears,_Waffles,_Data Science]"
```

```
pattern = r"[A-Z][a-z]+\W*\d{1,3}\w?"
```

List matches of the pattern above in the order they appear in the string, with the first match next to Match 1.

Note: You may have more or fewer than 6 total matches. For any unmatched slots, **write** No Match **instead**. If there are more than six matches, just write the first six that appear.

Match 1: '_____'

Match 2: '_____'

Match 3: '_____'

Match 4: '_____'

Match 5: '_____'

Match 6: '_____'

4 Sampling [4 Pts]

Sarah is the data scientist who oversees the **BayPass Program**. The BayPass Program is a study that analyzes the effects of giving UC Berkeley students free access to all Bay Area transportation services through a BayPass transit card.

- (a) Sarah takes a simple random sample of 100 UC Berkeley students from the UC Berkeley enrolled student database and gives them each a BayPass transit card. She collects usage data from these cards for analysis.

(i) [1 Pt] What is the population of interest of the study?

- ☐ All BayPass card holders
- ☐ All UC Berkeley students
- ☐ All Bay Area students who use Bay Area transit services
- ☐ All enrolled UC Berkeley students who use Bay Area transit services

(ii) [1 Pt] What is the sampling frame in Sarah's study? Be specific.

- (b) [2 Pts] Sarah considers using a stratified random sample to select 100 students, instead of a simple random sample. Why might Sarah want to consider a stratified random sample?

- ☐ True ☐ False She wants to guarantee that small but important groups are represented in the sample.
- ☐ True ☐ False She wants to reduce selection bias.
- ☐ True ☐ False She wants to reduce reporting bias (i.e., response bias).
- ☐ True ☐ False She wants to reduce chance error.
- ☐ True ☐ False She wants to reduce non-response bias.

5 Shell We Do Some Regression? [6 Pts]

Zara is building a model to predict how long it takes a turtle to finish a race, `finish_time` (in minutes), using features like `aura` (a turtle's charisma score) and `shell_diameter` (in cm).

- (a) [1.5 Pts] We now use the following model equation to predict `finish_time`:

$$\text{finish_time} = \hat{\theta}_0 + \hat{\theta}_1 \times \text{aura}$$

For each action, determine whether it would change the $\hat{\theta}_1$ value.

- ☐ True ☐ False Adding a constant to all values of `aura`.
☐ True ☐ False Multiplying all values of `aura` by 3.
☐ True ☐ False Using the square root of `aura` as the predictor variable.

- (b) [3 Pts] We now use the following model equation to predict `finish_time`:

$$\text{finish_time} = \hat{\theta}_0 + \sum_{j=1}^d \hat{\theta}_j \times \text{aura}^j$$

and train it with L2 regularization and hyperparameter λ . We perform 5-fold cross-validation and find the following table of cross-validation errors for each (λ, d) pair:

	$\lambda = 0.02$	$\lambda = 0.08$	$\lambda = 0.15$	$\lambda = 0.30$
$d = 1$	6.8	3.9	1.5	2.1
$d = 2$	5.9	0.7	2.4	4.0
$d = 3$	1.3	4.1	5.2	6.9

- (i) [1 Pts] How many times do we need to evaluate a model on a validation set to construct the above table?

- (ii) [1 Pts] Which pair of hyperparameters (λ, d) in the table above is expected to lead to the highest model bias squared?

- (iii) [1 Pts] Which pair of hyperparameters (λ, d) in the table above is expected to lead to the highest model variance?

- (c) [1.5 Pts] We now train the following model to predict `finish_time`:

$$\text{finish_time} = \hat{\theta}_0 + \hat{\theta}_1 \times (\text{shell_diameter})^2$$

We find the following optimal parameters for the model: $\hat{\theta}_0 = 2$, $\hat{\theta}_1 = 2$. We evaluate our model on the following dataset:

	shell_diameter	finish_time
Michaelangelo	1	1
Raphael	2	9
Donatello	3	13

What is the empirical risk if we fit the model with **L1 loss** and **L1 regularization** with hyperparameter $\lambda = \frac{2}{3}$? The intercept is not regularized. You can leave your answer as an unsimplified algebraic expression with numbers plugged in.

6 Spiderman: Far From the Minimum [13 Pts]

Kevin wants to estimate how long it takes Spiderman to get to an emergency, `response_time` (in minutes), using a dataset of Spiderman's past emergency responses, `spidey_calls`.

- (a) [2 Pts] Kevin is interested in creating a constant model ($\hat{y} = \theta_0$) to make predictions. He first must decide between optimizing MSE or MAE. Select all statements below that are true.

- ☐ It is possible for two different constant predictions to both result in the optimal MAE.
- ☐ Since the MSE loss surface is always convex and smooth, there is always one optimal θ_0 that minimizes it.
- ☐ Kevin can use gradient descent to find a value of θ_0 that minimizes MAE.
- ☐ Suppose Kevin chooses to optimize the MSE. The median of the data could result in a lower training loss than the training loss associated with the mean.

- (b) [1 Pt] Suppose Kevin's design matrix $\mathbb{X}$ has full column rank, so he generates unique OLS predictions $\hat{\mathbb{Y}}$. Which of the following is true?

- ☐ True ☐ False The residual vector $\mathbb{Y} - \hat{\mathbb{Y}}$ has the same dimension as the parameter vector $\vec{\theta}$.
- ☐ True ☐ False The predictions $\hat{\mathbb{Y}}$ are the projection of $\mathbb{Y}$ onto the span of $\mathbb{X}$.
- ☐ True ☐ False If $\mathbb{X}$ is full column rank, there is exactly one $\hat{\theta}$ vector that minimizes the mean squared error.
- ☐ True ☐ False The prediction vector $\hat{\mathbb{Y}}$ is a linear combination of the columns of $\mathbb{X}$.

- (c) [2 Pts] Kevin first fits the following simple model to predict `response_time` using only `blocks_away`:

$$\text{response_time} = \hat{\theta}_0 + \hat{\theta}_1 \times (\text{blocks_away})$$

- (i) [1 Pt] In one sentence, interpret the meaning of $\hat{\theta}_0$ in this model.

- (ii) [1 Pt] Kevin now one-hot encodes the categorical feature `borough` for which NYC borough Spiderman is currently in (which takes on values `Manhattan`, `Brooklyn`, or `Queens`), and adds two new features to his model, leaving out the `Manhattan` column:

$$\text{response_time} = \hat{\theta}_0 + \hat{\theta}_1 \times (\text{blocks_away}) + \hat{\theta}_2 \times (\text{Brooklyn}) + \hat{\theta}_3 \times (\text{Queens})$$

In one sentence, interpret the meaning of $\hat{\theta}_0$ in this **new** model.

- (d) [4 Pts] Kevin then tries a completely different loss function to predict `response_time` using two features:

- x_{i1} : blocks away from the emergency on call i
- x_{i2} : wind speed on call i

He defines the following regularized loss function:

$$L(\theta_1, \theta_2) = \frac{1}{2n} \sum_{i=1}^n (y_i - \theta_1 x_{i1} - \theta_2 x_{i2})^2 + \frac{\lambda}{2} (\theta_1^2 + \theta_2^2).$$

Kevin decides to minimize this loss using gradient descent.

- (i) [2 Pts] Derive $\frac{\partial L}{\partial \theta_1}$. Show your work.

He starts from $\theta_1^{(0)} = 0$, $\theta_2^{(0)} = 0$, uses a learning rate $\alpha = 0.1$, and a regularization term $\lambda = 0.5$. **Kevin uses new data to fit this model:**

$$\mathbb{X} = \begin{bmatrix} 1 & 3 \\ 2 & 1 \\ 4 & 2 \end{bmatrix}, \quad \mathbb{Y} = \begin{bmatrix} 10 \\ 15 \\ 20 \end{bmatrix}$$

(ii) [2 Pts] Use gradient descent to calculate $\theta_1^{(1)}$. Show your work.

(e) [2 Pts] Three different training paradigms are used to optimize a model. Each runs for **5 epochs**.

Paradigm	Number of Data Points
Paradigm 1 (Batch GD)	150
Paradigm 2 (Mini-batch GD)	200
Paradigm 3 (Stochastic GD)	10

Which ordering correctly ranks the paradigms by total number of gradient computations **across all epochs**?

Note: Paradigm 2 uses a batch size of 10.

- ☐ Paradigm 1 < Paradigm 2 < Paradigm 3
- ☐ Paradigm 1 < Paradigm 3 < Paradigm 2
- ☐ Paradigm 3 < Paradigm 1 < Paradigm 2
- ☐ Paradigm 2 < Paradigm 3 < Paradigm 1

(f) [2 Pts] Select all statements below that are true regarding gradient descent.

- ☐ With an **infinite** number of updates and a **fixed** learning rate, gradient descent will always converge to a **global** minimum.
- ☐ With an **infinite** number of updates and a **fixed** learning rate, gradient descent will always converge to a **local** minimum.
- ☐ The gradient descent algorithm can compute gradient updates using both convex and non-convex functions.
- ☐ Consider the average of the estimated gradients computed over one epoch of stochastic gradient descent. This value is always equal to the true gradient of the loss surface at the initial values of the model parameters.

7 Hit Me With Your Best Shot [5 Pts]

Rohan wants to model his rate of perfect shots using the estimator:

$$\hat{p} = \frac{1}{n} \sum_{i=1}^n X_i$$

Let X_i be an indicator random variable that equals 1 if Rohan makes a perfect shot on trial i , and 0 otherwise, for $i = 1, \dots, n$. Each trial is independent, so $X_1, \dots, X_n$ are i.i.d. $\text{Bernoulli}(p)$ random variables. Let the actual probability of Rohan making a perfect shot be an unknown constant p (i.e., $P(X_i = 1) = p$ for each i). For each subpart below, grading will be done based on the work you show in the boxes below.

- (a) [2.5 Pts] What is the variance of $\hat{p}$? **Express and simplify your answer in terms of n , $\hat{p}$, p , and any real number.**

- (b) [2.5 Pts] What is the bias of this estimator: $\mathbb{E}[\hat{p}] - p$? **Express and simplify your answer in terms of n , $\hat{p}$, p , and any real number.**

8 A Messi Affair [9 Pts]

The 2026 World Cup is underway, and Aadil wants to predict whether Messi's Argentina wins their next match. Let $y_i = 1$ if Messi wins match i , and $y_i = 0$ otherwise.

- (a) [3.5 Pts] Aadil wants to build a logistic regression model using design matrix $\mathbb{X}$ to predict whether Argentina wins. This matrix has two features: $\mathbb{X}_{[:,1]}$ (`kids_there`, a boolean indicator of whether one of Messi's kids is at the game) and $\mathbb{X}_{[:,2]}$ (`rest_days`, the number of rest days before the match).

He builds a logistic regression model using $\mathbb{X}_{[:,1]}$ and $\mathbb{X}_{[:,2]}$ to predict whether Argentina wins ($y_i = 1$) or loses ($y_i = 0$). The model uses the following parameter vector:

$$\vec{\theta} = [-2 \quad 2 \quad 1]^\top$$

and follows the standard logistic regression form.

For a particular match m , one of Messi's kids was there and the team had 0 rest days before the match.

- (i) [1.5 Pts] Compute the probability $\hat{p}$ that Argentina wins match m .

- (ii) [.5 Pts] Using a probability threshold of 0.45, does the model predict Argentina wins or loses match m ?

- ☐ Argentina loses
- ☐ Argentina wins
- ☐ On the boundary
- ☐ Not enough information

- (iii) [1.5 Pts] Derive the equation of the decision boundary using a probability threshold of 0.5. Express your answer in terms of $\mathbb{X}_{[:,1]}$ and $\mathbb{X}_{[:,2]}$. *Note: You may rearrange the equation so that constants are on either side. For example, $a^2 + b = 5$ would receive the same credit as $a^2 + b - 5 = 0$.*

- (b) [1 Pt] Evaluate the following statements about logistic regression models.

- ☐ True ☐ False For any $t \geq 0$, the sigmoid function σ satisfies $\sigma(t) \in [0.5, 1)$.
- ☐ True ☐ False In binary classification, a logistic regression model can predict the probability of a datapoint belonging to class 0.
- ☐ True ☐ False For all linearly separable datasets, there exists a finite number of optimal solutions for model parameters.
- ☐ True ☐ False Let $\mathbb{X}$ be the design matrix. The quantity $\mathbb{X}^\top \mathbb{X}$ must be invertible in order to compute the optimal parameter vector.
- ☐ True ☐ False It is possible for the sigmoid function σ to output exactly $P(y_i = 1|x_i) = 1$ for some datapoint x_i with real-valued features.

- (c) To evaluate Aadil's logistic regression model, he tests it on 6 new matches and compares predicted versus true outcomes. Results are shown below:

Match	y	$\hat{y}$	$\hat{P}_\theta(y = 1 x)$
1	1	1	0.75
2	1	1	0.65
3	0	1	0.55
4	1	0	0.48
5	1	0	0.35
6	0	0	0.20

As a reminder, $y = 1$ indicates Argentina won the match, and $\hat{y}$ is the model's predicted label at whatever threshold it currently uses.

- (i) [1 pt] What is the recall of the model on this data? Leave your answer as a fraction.


- (ii) [1 pt] Given the values in the table above, what are the minimum and maximum possible values of the threshold used to produce these predictions? *For full credit, your answer does not need to address whether the bounds are inclusive or exclusive.*

- (iii) [2.5 pts] What will happen to this logistic regression classifier's metrics if the classifier's threshold is lowered? Evaluate the following statements.

- ☐ True ☐ False Precision will never increase.
☐ True ☐ False Recall will never decrease.
☐ True ☐ False It is possible that the model will predict "Messi wins" more often.
☐ True ☐ False It is possible that the model's ROC-AUC will change.

9 [12 Pts] Dimensionality, Reducto!

In the parallel wizarding world, Professor Josh is looking at the performance of **100 students** in his TaDa 100 class, he records each student's **Charms Score** (x_1) and **Potions Score** (x_2) (**both mean-centered**). Below are six different possible versions of this dataset, labeled Dataset A through Dataset F. Each scatterplot is centered at the origin and drawn on the same axis scale.



(a) [2 Pts] Josh computes the covariance matrix of one of the six datasets above to be:

$$\Sigma = \begin{bmatrix} 4 & 0 \\ 0 & 36 \end{bmatrix}$$

Which scatterplot could this covariance matrix correspond to?

- ☐ Dataset A
- ☐ Dataset B
- ☐ Dataset C
- ☐ Dataset D
- ☐ Dataset E
- ☐ Dataset F

(b) [2 Pts] Consider this ratio:

$$\frac{\text{variance of the 2nd latent feature}}{\text{variance of the 1st latent feature}}$$

For which of the following datasets is this ratio the **largest**?

- ☐ Dataset A
- ☐ Dataset B
- ☐ Dataset C
- ☐ Dataset D
- ☐ Dataset E
- ☐ Dataset F

(c) [2 Pts] Josh takes one of the six datasets above, **samples 5** representative points, and computes the following (approximate) SVD of the resulting design matrix $\mathbb{X}$.

$$U \approx \begin{bmatrix} -0.71 & 0.10 \\ 0.24 & -0.55 \\ 0.47 & 0.60 \\ -0.35 & 0.30 \\ 0.35 & -0.40 \end{bmatrix} \quad S \approx \begin{bmatrix} 20 & 0 \\ 0 & 3 \end{bmatrix} \quad V^T \approx \begin{bmatrix} 0.71 & 0.71 \\ -0.71 & 0.71 \end{bmatrix}$$


(i) Which scatterplot could this SVD correspond to?

- ☐ Dataset A
- ☐ Dataset B
- ☐ Dataset C
- ☐ Dataset D
- ☐ Dataset E
- ☐ Dataset F

- (d) [2 Pts] Using the SVD given in part (c), what is the **component score** of the first principal component? In other words, what is the variance of the five sampled datapoints in the direction of the first principal component? You can leave your answer as an unsimplified algebraic expression (e.g., $(5 + 3)/2 + 1$). The matrix is repeated below for your convenience:

$$U \approx \begin{bmatrix} -0.71 & 0.10 \\ 0.24 & -0.55 \\ 0.47 & 0.60 \\ -0.35 & 0.30 \\ 0.35 & -0.40 \end{bmatrix} \quad S \approx \begin{bmatrix} 20 & 0 \\ 0 & 3 \end{bmatrix} \quad V^T \approx \begin{bmatrix} 0.71 & 0.71 \\ -0.71 & 0.71 \end{bmatrix}$$

(e) [2 Pts] Cole projects a different sampling of points onto the first principal component.



Cole decides to use **K-medians clustering**. **K-medians clustering is exactly the same as K-means clustering, but it uses the median instead of the mean as the cluster centers.** He initializes Cluster 1's center at 0 and Cluster 2's center at 7.

Write the cluster center updates of K-medians clustering for this dataset. Write the updates in order and **until the algorithm converges**. We have provided a template for you to fill in. Each row represents one iteration of K-medians clustering.

Write the location of the updated Cluster 1 center and Cluster 2 center after each assignment step. **You do not have to use all rows.** If the algorithm has already converged (i.e., the centers don't change), please write NA in any boxes you do not use.

*Note: Remember Cole is using **K-medians clustering**, not K-means clustering. If you need to find the median of an even number of points, use the average of the two middle points.*

Step	Cluster 1 center	Cluster 2 center
Initialization	0	7
Step 1	_____	_____
Step 2	_____	_____
Step 3	_____	_____
Step 4	_____	_____

(f) [2 Pts] Which of the following statements are true?

- ☐ True ☐ False A negative silhouette score suggests that a point is, on average, closer to the points in the nearest neighboring cluster compared to the points in its own cluster.
- ☐ True ☐ False Agglomerative clustering will always converge to the same solution if re-run with the same parameters. *For this option, you can assume there are never exact ties in distances between clusters.*
- ☐ True ☐ False For a fixed dataset and a fixed value of K , K-means clustering always produces the same final cluster assignments, regardless of the initial placement of centroids.
- ☐ True ☐ False K-means optimizes the silhouette score.

Please state any relevant assumptions in the box below (Optional).

You are done with the final- Congratulations!

Draw your favorite DATA 100 memory so far!

A large, empty rectangular box with a thin black border, intended for a drawing. It occupies the majority of the lower half of the page.